

THE NARRATIVE

Strategy, Product & Execution Plan

Confidential · generated from the approved plan

The Narrative - Strategy + Product + Execution (OSINT-led)

*Reframes the project as an **OSINT platform** whose differentiator is a consequence engine.*

*This supersedes both the OSINT-only plan and the consequence-only framing: the consequence work stays - it is what makes our OSINT one-of-a-kind. Identity = **Signal -> Consequence**.*

*The 3 founder forks (market, GTM, openness) are **decided** - see §11.*

1. Core idea (as it stands now)

The Narrative = an open-source-intelligence surface that carries a signal to its consequence.

It is one place to investigate any entity - IP, domain, actor, wallet, vessel, claim - across

1,098 OSINT tools with **live keyless enrichment**, and then it does the thing every other

OSINT product stops short of: it maps that signal to its **2nd/3rd-order consequences** (supply chain -> sector -> region -> citizen/asset impact), shows **evidence at every node**, makes **calibrated predictions**, and **grades its own accuracy** over time.

It runs **\$0 / local-first** (free authoritative feeds + the OSINT catalog, local LLM, deterministic consequence engine), with paid AI as an opt-in, hard-capped escalation - making it **private / on-prem-capable** by construction.

One line: *****Most OSINT stops at the link. We carry the signal to consequence - scored and self-graded, at near-zero marginal cost.*****

2. The real-world problem we solve

OSINT today is a **directory of tools and a pile of links**. An analyst pivots an indicator across dozens of sites by hand, copies values between tabs, and is then **left at the link** - with no answer to "so what does this *do* to us?" Incumbent intelligence platforms go the other way: they surface alerts but (a) rarely reason about **cascading impact** with auditable evidence, (b) **never grade their own forecast accuracy** (trust is faith-based), (c) are ****expensive and cloud-only**** - a non-starter for privacy/regulated buyers whose data can't leave.

We close both gaps at once: a **one-click OSINT investigation surface** and automated, evidence-backed **consequence mapping** with a track record you can check.

3. The product, concretely (what makes our OSINT one-of-a-kind)

- **The catalog as actions, not a directory.** All 1,098 vendored OSINT-Framework tools are classified by what they can *actually do*, honestly, in three tiers:
 - **live** - keyless server-side enricher returns real facts in-app (highest capability).
 - **pivot** - the tool takes an entity, so we pre-fill it: into the tool's **native** search URL (curated, host-keyed) or, for the long tail, an honest **site-scoped** search. One click either way.
 - **launch** - no entity to pre-fill, or a desktop install: we open it and copy the value.
 - A curated **entityKinds override** recovers tools the vendored snapshot mis-tagged as inert (NVD -> CVE, Shodan -> IP/domain, VIN decoders -> vehicle, hash DBs, XMRChain, HLR) ****without** mutating the vendored data. **Guarantee:** 0 dead links** across all 1,098 tools.
- **Per-category entity investigators** + a **keyless threat/disinfo feed** sit on top of the catalog.
- **The consequence engine** is the differentiator - the auditable causal chain from any investigated signal to sector/region/citizen impact, with calibrated, self-graded predictions.

4. Who needs it - ICP (all segments; land-and-expand)

1. **SOC / threat-intel / DFIR teams** - the native OSINT audience: pivot an indicator, then see its downstream consequence. Fastest credibility for the OSINT framing.
2. **Enterprise risk / resilience / supply-chain & security teams** - "what does this event do to *our* operations, suppliers, sites?" Clear budget, clear pain.
3. **Finance - funds, commodities, macro desks** - geopolitics/indicators -> market/sector consequence.
4. **Government / defense / NGOs** - situational awareness **with on-prem/privacy** (our wedge).
5. **Prosumers - analysts, journalists, researchers** - cheap self-serve tier; bottom-up adoption.

Decision (you chose "all"): all are the target market (full TAM). We **land-and-expand** - ship one product, turn on segments in sequence: **(1) analyst/SOC + prosumer self-serve** (the OSINT crowd validates the surface and feeds the flywheel) -> **(2) enterprise risk/supply-chain** (first real ACV) -> **(3) finance** -> **(4) gov/defense on-prem** (highest margin) -> **(5) insurers**. Same core product throughout - only packaging and messaging change per segment.

5. Competition & how we beat them

Competitor	Strength	Their gap = our wedge
OSINT Framework / Maltego / SpiderFoot	Breadth of tools / link directory / pivoting	Stops at the link - no consequence, no scoring, no self-grading
Seerist	AI threat/geopolitical risk + analyst network	Surfaces risk, not downstream consequence ; no self-grading; \$\$\$ cloud-only
World Monitor	Real-time, map-centric event monitoring	Event-level only; no consequence reasoning/calibration

Competitor	Strength	Their gap = our wedge
DataSurfr	Data feeds + dashboards/API	Data plumbing, not interpretation; no impact graph
Recorded Future / Dataminr / Factal	Breadth, speed, brand	Very expensive, cloud-only, opaque scoring

****We don't out-breadth the OSINT directories - we make their tools *actionable* and then out-*reason* and out-*price* the intelligence incumbents**, and we serve buyers they structurally can't (privacy/on-prem). Four levers:**

1. **OSINT that finishes the job** - every investigation carries to an auditable consequence chain.
2. **Calibrated, self-grading predictions** - published accuracy -> trust + a proprietary dataset.
3. **\$0 / local-first + on-prem** - undercut on price *and* win privacy-bound segments.
4. **Owned trained models** - near-zero marginal inference cost => structurally cheaper than rivals.

6. Moat / defensibility (the flywheel)

- **Proprietary calibrated-outcome dataset** - every prediction + its graded result. Compounds with usage; cannot be bought.
- **Owned fine-tuned models** (triage SLM -> mapper SLM) trained on that data -> \$0 marginal cost, on-prem-deployable. (See §8.)
- **Auditable consequence engine + evidence graph** as defensible IP.
- **The curated capability/override layer** over the OSINT catalog - the honest "what each tool can do" map - is itself a maintained asset rivals' static directories don't have.
- **Flywheel:** usage -> more graded outcomes -> better models/calibration -> better product -> usage.

7. Business model & GTM (decisions locked)

- **Closed-source / proprietary** (your call) - engine, trained models, capability layer, and calibrated dataset all private. We earn trust via the **public accuracy track record** (§6) rather than open code. (The underlying OSINT *tool list* is open; our *actioning* + *consequence* of it is the product.)
- **Hybrid GTM, sequenced** (you chose "all"): **SaaS self-serve tiers first** (free -> Pro -> Intelligence/Enterprise; already in `web/src/lib/tiers.js`) to get the OSINT crowd and flywheel data fast, then **on-prem/private** as the enterprise + gov/defense upsell.
- **Motion:** per land-and-expand (§4) - 1-2 design partners in the active segment, a concrete "indicator -> consequence" win, then expand.

8. Model strategy - local LLM vs trained SLM vs API

****Recommendation:** hybrid - local-first now -> distilled SLM next -> paid API only as a capped escalation.** Scored for *this* project (high-volume triage/mapping, \$0 posture, privacy, moat):

Option	Fit	Cost	When
Local LLM (Ollama, 3-8B)	Good for triage/extraction	\$0	now / default hot path
Trained SLM (fine-tuned)	Best long-term - fast, consistent, moat	\$0 serve; ~\$2-15 train	next , once enough labeled data
Paid API (Claude)	Best raw quality	per-call \$\$; data egress	hard cases + label generation , hard-capped

Why: triage/mapping run at volume -> cost & latency dominate. The **data flywheel** (\$6) is the training set; **triage is the best first distillation target** (narrow, high-volume, trivially evaluable). No re-architecture: backend/services/llm.py already abstracts provider; add (a) decision logging, (b) later an slm provider via Ollama. ****API never the default - it breaks \$0 and ships data out.****

9. Product roadmap (phased)

Now - OSINT surface is live; harden sourcing

- All 1,098 catalog tools actioned with honest capability tiering + the entityKinds override (0 dead links). [x] shipped.
- OSINT v2 **multi-source collector** (one failing source never kills ingestion): RSS/Atom (incl. Reddit .rss, Google News), then GDELT (keyless DOC API), then Reddit OAuth + Mastodon. All flow through `osint_agent.triage()` -> `_upsert`.
- **Turn the scheduler on + rotate the leaked Anthropic/Voyage keys.**

Next - the differentiator

- **Corroboration/verification agent**: cluster OSINT vs authoritative feeds by embedding+geo+time; ≥ 2 independent sources => promote OSINT -> **verified**. Feeds the flywheel.
- Wire investigated entities -> consequence chains end-to-end (the "carry to consequence" promise).

Then - moat + parity

- **Distill the triage SLM** from logged data; later the mapper SLM.
- Native search-pattern coverage sweep (raise pivots from site-scoped -> native), exports/API surface.

10. Tech posture (keep)

FastAPI + pgvector + async worker scheduler + React/d3 globe; deterministic **consequence engine** for auditable chains; the **OSINT capability/override layer** over the vendored catalog; **provider-abstracted** inference with enforced `cost_guard` caps; \$0/local-first default; pluggable free-feed "Signal" pattern. Evolve, don't rewrite.

11. Risks & how we tackle them

- **Source/tool blocking** -> multi-source portfolio + honest tiering (site-scoped fallback never dies).
- **Trust in AI claims** -> evidence-per-node + self-grading + corroboration.

- **Incumbent breadth** -> compete on actionable OSINT + consequence + on-prem, not raw breadth.
- **Local-model quality** -> distillation + API escalation for hard cases only.
- **Leaked keys / data egress** -> rotate keys; \$0-local + on-prem keeps data in-house.

12. Decisions locked (founder)

1. **Identity = OSINT-led, Signal -> Consequence.** Lead the brand/IA with OSINT; the consequence engine is the one-of-a-kind differentiator (kept, not reverted).
2. **Market = all segments**, land-and-expand (§4): analyst/SOC -> enterprise risk -> finance -> gov/defense on-prem -> insurers.
3. **GTM = hybrid, sequenced** (§7): SaaS self-serve first, on-prem upsell second.
4. **Closed-source / proprietary** - moat = data + trained models + capability layer + engine; trust comes from the published accuracy record (§6).

13. Immediate next steps

1. Verify the live stack: rotate keys -> start scheduler -> run OSINT once.
2. Wire investigated entities -> consequence chains (deliver the headline promise end-to-end).
3. Stand up **triage-decision logging** (the flywheel) - cheap, compounds early.

Verification (engineering, per increment)

- `python -m backend.services.osint_catalog_test green` (capability tiering + 0-dead-link reachability).
- `python -m backend.feeds.feeds_test green` (pure parser tests per source).
- `python -m backend.workers.osint_ingest_worker live` -> **non-zero** posts/new; rerun => new:0 (dedup).
- Corroboration: two near-dupe items from different sources => one promoted to "verified by N".
- Degrade check: stop Ollama => heuristic triage still ingests.